

# Machine Learning B (2025)

## Home Assignment 5

Carsten Jørgensen, student ID: skj730

### Contents

|   |                          |   |
|---|--------------------------|---|
| 2 | PAC-Bayesian Aggregation | 2 |
| 3 | VC Dimension             | 6 |

## 2 PAC-Bayesian Aggregation

The main task of exercise 3.8 from the Lecture Notes [Sel25] is to reproduce Figure 2(a) from [Thi+17].

I am using the SVM algorithm from [Ped+11] for the classification task. This algorithm is a wrapper around LIBSVM [CL11], which is being used in [Thi+17].

Here is my version of Figure 2(a) 1.




Figure 1: Comparison of PAC-Bayesian aggregation with RBF kernel SVM tuned by cross-validation.

### Data Preparation

The data set used for Figure 2(a) is the *Ionosphere* data set from [Sig+89]. The original data has class labels  $b$  and  $g$  remapped to  $-1, 1$ . Label  $b$  is bad and  $g$  is good.

The data is then split into a training set  $S$  and test set  $T$  with  $|S| = 200$ ,  $|T| = 150$  as mentioned in Table 1 in [Thi+17]<sup>1</sup>. This split was creating by scikit-learn

<sup>1</sup>The data set actually contains 351 rows so I am just ignoring a random point to match size of training and test sets given by [Thi+17].

method `train_test_split` with a stratify approach to ensure balance of labels.

All features have been standardized. I cannot find any where in [Thi+17] mentioning this, but standardization of features is recommended in Section 2.2 of the LIBSVM tutorial [HCL03].

## Baseline cross-validated SVM

Paragraph 2 of Section 6 in [Thi+17] explains how the Baseline cross-validated SVM is constructed using a grid search with parameters  $C$  and  $\gamma$ .

The  $\gamma$ -grid is based on the Jaakkola heuristic from [JDH99] and the calculation of the  $\gamma$  value is basically the median distance between pairs of points with opposite labels. My implementation of the heuristic is a wrapper around the `cdist` function from [Ped+11].

Using `GridSearchCV` the best parameters across searched parameters are

$$C = 10.0, \quad \gamma = 0.0071686$$

yielding a model score of 0.945 and one-zero loss of 0.086092

## Alternating minimization of the PAC-Bayes- $\lambda$ bound

This is the core part of being able to reproduce Figure 2(a). The complete code listing for the Python implementation of the alternating minimization is listed in Section 3. Here we shall only include the most essential parts.

The input to the algorithm are the

**Validation losses (Eq. (13)).** For every weak learner  $h$  we need the empirical loss  $\hat{L}_{\text{val}}(h, S)$  on *validation* points  $S \setminus S_h$ . This is done inside the training loop, see Section 3 for the `L_val` variable.

**Numerically stability of the update rule for  $\rho$  in Eq.(7).** The exercise hints a way to ensure a numerical stable calculation of the Eq.(7). Here I am using the log-sum-exp trick as explained in [Gun99] and if I understand the hint correctly, then I believe that it is more or less equivalent to the log-sum-exp trick. This is implemented as:

```
# Eq. (7): computed in log-domain for numerical stability
log_pi = -np.log(m) * np.ones(m) # uniform prior \pi(h)=1/m (p.
    10)
log_rho = log_pi - lambda_ * n_r * L_val
log_rho -= logsumexp(log_rho)
rho = np.exp(log_rho)
```

**KL divergence and empirical loss.** To compute the KL term, I am using the first equality from Eq. (9):

$$KL(\rho||\pi) = \mathbb{E}_{h \sim \rho} \left[ \ln \frac{\rho(h)}{\pi(h)} \right] = \sum_{h \in \mathcal{H}} \rho(h) \ln \frac{\rho(h)}{\pi(h)}.$$

to keep the calculation numerical stable in the log-domain.

```
# helpers for Eq. (8) & bound. KL divergence is computed
# in log-domain using the first equality in Eq. (9)
KL = np.sum(rho * (log_rho - log_pi))
L_emp = np.dot(rho, L_val)
```

**Closed-form update of  $\lambda$  from Eq.(8).** This is the code for updating  $\lambda$ :

```
# Eq. (8): lambda_new
num = 2.0
denom = (
    np.sqrt(2.0 * n_r * L_emp / (KL + np.log(2 * np.sqrt(n_r) /
        delta)))
    + 1.0
    + 1.0
)
lambda_new = num / denom
```

from Eq.(8):

$$\lambda^{\text{new}} = \frac{2}{\sqrt{\frac{2(n-r)\mathbb{E}_\rho[\widehat{L}(h, S)]}{KL(\rho||\pi) + \ln \frac{2\sqrt{n-r}}{\delta}} + 1 + 1}}.$$

**Upper Bound value (right-hand side of Eq.(13)).**

$$F(\rho, \lambda) = \frac{E_\rho[\widehat{L}]}{1 - \frac{\lambda}{2}} + \frac{KL(\rho||\pi) + \ln \frac{2\sqrt{n-r}}{\delta}}{\lambda(1 - \frac{\lambda}{2})(n-r)}.$$

```
# Eq. (13): value of PAC-Bayes-\lambda bound
bound = L_emp / (1 - lambda_ / 2.0) + (KL + np.log(2 * np.sqrt(n_r)
    ) / delta)) / (
    n_r * lambda_ * (1 - lambda_ / 2.0)
)
```

## Construction of Hypothesis Space

For each  $m^2$  on the  $x$ -axis of Figure 2(a) we create  $m$  weak-learners SVM and compute the  $m$  empirical losses.

---

<sup>2</sup>The paper does not specify the  $m$ -value but I have eye-balled them to be 1, 2, 3, 4, 5, 6, 8, 10, 12, 15, 20, 25, 35, 45, 65, 85, 110, 135, 160, 200,

## Majority vote evaluation

For the majority vote I use Eq. 3.31 in [Sel25]:  $MV_{\rho}(X) = \text{sign}\left(\sum_{h \in \mathcal{H}} \rho(h)h(X)\right)$  implemented in the last lines of code in the training algorithm 2.

## Conclusion

The main difference that I see between Figure 2(a) from [Thi+17] and my reproduction seems to be the running time for the baseline cross-validated SVM. I am using the scikit-learn SVM which is a wrapper around LIBSVM used in the paper. Maybe that can explain some of the performance loss. In addition the paper does not specify what kind of hardware that was used to produce the results.

In Figure 2 we see an updated version of Figure 2(a) where the experiment was repeated 10 times such that a 95% confidence level could be computed for the bounds.

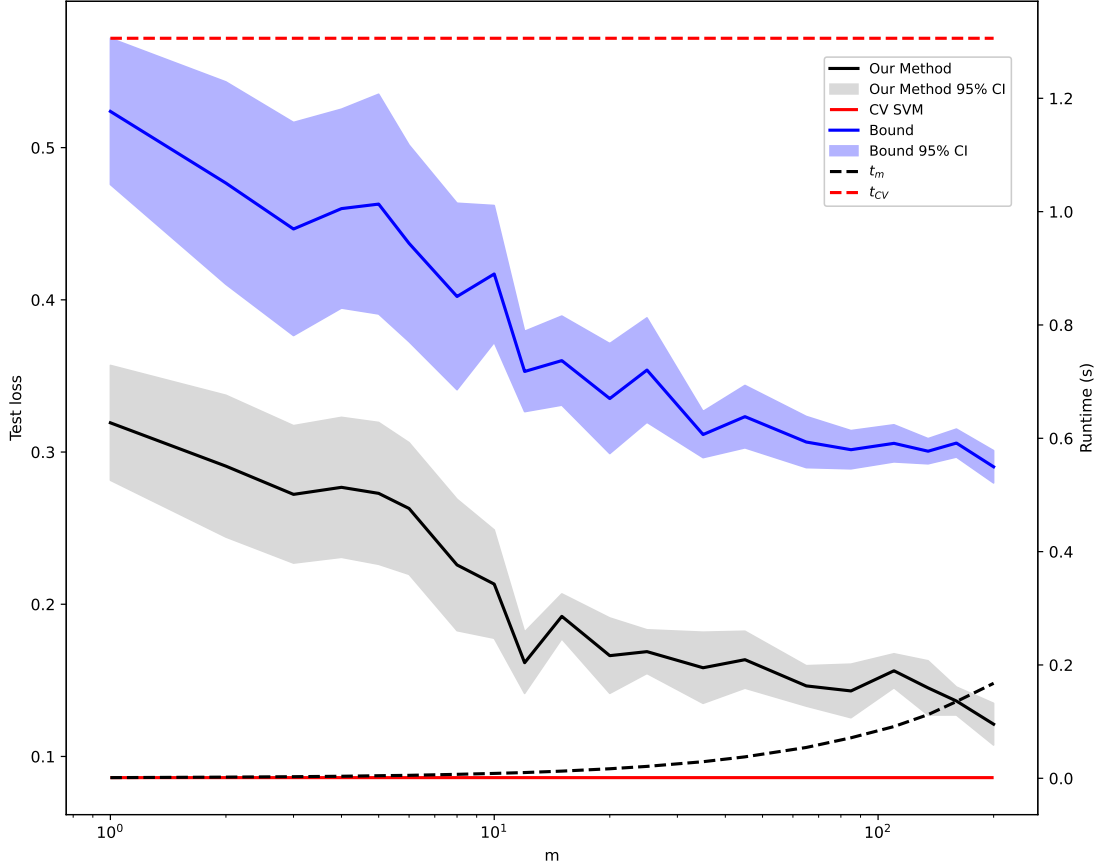


Figure 2: Figure 2(a) with 95% Confidence intervals (10 runs). The solid blue and black line are the means of 10 experiments for each  $m$ . The confidence intervals represent  $\pm 1.96SE$  where  $SE$  is the standard error on the mean. Note running times are plotted for a single run.

Table 1 summarizes the statistics for 10 runs of the experiment. The details of the

baseline cross-validated SVM model are: cross-validation score: 0.9450 and best parameters:  $C = 10.0$ ,  $\gamma = 0.007065$ .

| Method          | Loss ( $\mu \pm \sigma$ ) | Runtime ( $\mu \pm \sigma$ ) [s] |
|-----------------|---------------------------|----------------------------------|
| Our Method      | $0.1212 \pm 0.0222$       | $0.1904 \pm 0.0069$              |
| PAC-Bayes Bound | $0.2903 \pm 0.0172$       | -                                |
| CV SVM          | 0.0861                    | 1.2932                           |

Table 1: Statistics for 10 repeated experiment

### 3 VC Dimension

#### Q1 of Exercise 3.4

##### Finite hypothesis class

*Proof.* Let  $\mathcal{H}$  be finite,  $|\mathcal{H}| = M$ . If a set of  $n$  points is shattered by  $\mathcal{H}$ , then

$$2^n \leq |\mathcal{H}| = M.$$

Hence  $n \leq \lfloor \log_2 M \rfloor$  and therefore

$$d_{VC}(\mathcal{H}) \leq \lfloor \log_2 M \rfloor.$$

□

#### Q2 of Exercise 3.4

**Two-element hypothesis class** Let  $\mathcal{H}$  be a hypothesis space with 2 hypotheses (i.e.,  $|\mathcal{H}| = 2$ ). Prove that  $d_{VC}(\mathcal{H}) = 1$ .

*Proof.* To prove equality, we need to show both  $d_{VC}(\mathcal{H}) \geq 1$  and  $d_{VC}(\mathcal{H}) \leq 1$ .

**Upper bound:** From Problem 1, we have:

$$d_{VC}(\mathcal{H}) \leq \lfloor \log_2(2) \rfloor = 1$$

**Lower bound:** Let  $\mathcal{H} = \{h_1, h_2\}$  where  $h_1 \neq h_2$ . Since the hypotheses are different, there exists at least one point  $x$  where  $h_1(x) \neq h_2(x)$ . Without loss of generality, assume  $h_1(x) = +1$  and  $h_2(x) = -1$ . For the single point  $\{x\}$ , we can achieve both possible labelings:

- Labeling (+1): use hypothesis  $h_1$
- Labeling (-1): use hypothesis  $h_2$

Since we can achieve all  $2^1 = 2$  possible labelings of one point, the set  $\{x\}$  can be shattered. Therefore,  $d_{VC}(\mathcal{H}) \geq 1$ . Combining both bounds:  $d_{VC}(\mathcal{H}) = 1$ . □

### Q3 of Exercise 3.4

**Positive circles in  $\mathbb{R}^2$**  Let  $\mathcal{H}_+$  be the class of “positive” circles in  $\mathbb{R}^2$  (each  $h \in \mathcal{H}_+$  is defined by the center of the circle  $c \in \mathbb{R}^2$  and its radius  $r \in \mathbb{R}$ ; all points inside the circle are labeled positively and outside negatively). Prove that  $d_{VC}(\mathcal{H}_+) \geq 3$ .

*Proof.* We need to show that there exists a set of 3 points that can be shattered by positive circles.

Consider the following three points that are not collinear:

$$A = (0, 0)$$

$$B = (1, 0)$$

$$C = (0, 1)$$

We will show that all  $2^3 = 8$  possible labelings can be achieved:

1.  $(-, -, -)$ : Use a small circle far from all points
2.  $(+, -, -)$ : Small circle centered at  $A$  with radius  $< 1$
3.  $(-, +, -)$ : Small circle centered at  $B$  with radius  $< 1$
4.  $(-, -, +)$ : Small circle centered at  $C$  with radius  $< 1$
5.  $(+, +, -)$ : Circle that includes  $A$  and  $B$  but not  $C$  (e.g., center around  $(0.5, -0.5)$ )
6.  $(+, -, +)$ : Circle that includes  $A$  and  $C$  but not  $B$  (e.g., center around  $(-0.5, 0.5)$ )
7.  $(-, +, +)$ : Circle that includes  $B$  and  $C$  but not  $A$  (e.g., center around  $(1.5, 1.5)$ )
8.  $(+, +, +)$ : Large circle containing all three points

Since all 8 labelings are achievable, these 3 points can be shattered by  $\mathcal{H}_+$ . Therefore,  $d_{VC}(\mathcal{H}_+) \geq 3$ .  $\square$

### Q4 of Exercise 3.4

**Positive & negative circles** Let  $\mathcal{H} = \mathcal{H}_+ \cup \mathcal{H}_-$  be the class of “positive” and “negative” circles in  $\mathbb{R}^2$  (the “negative” circles are negative inside and positive outside). Prove that  $d_{VC}(\mathcal{H}) \geq 4$ .

*Proof.* We need to show that there exists a set of 4 points that can be shattered by the union of positive and negative circles. Consider the following four points forming a unit square:

$$\begin{aligned} A &= (0, 0) \\ B &= (1, 0) \\ C &= (1, 1) \\ D &= (0, 1) \end{aligned}$$

For any subset  $S$  of the 4 points that we want to label positive:

- If  $S$  can be enclosed by a circle without including other points  $\Rightarrow$  use a positive circle
- If the complement of  $S$  can be enclosed by a circle without including points in  $S \Rightarrow$  use a negative circle

The 4 points in a square configuration have the property that for any subset, either the subset or its complement can be reasonably enclosed by a circle. Examples of achievable labelings:

- $(+, +, -, -)$  on  $A, B, C, D$ : Use negative circle around  $C, D$
- $(+, -, +, -)$  on  $A, B, C, D$ : Use negative circle around  $B, D$
- Other alternating patterns: Use appropriate negative circles

Since we can handle all  $2^4 = 16$  possible labelings using the combination of positive and negative circles, these 4 points can be shattered. Therefore,  $d_{VC}(\mathcal{H}) \geq 4$ .  $\square$

## Q5 - Exercise 3.6

**The fine details of the lower bound** Hypothesis Space  $\mathcal{H}$ :

$$\mathcal{H} = \{h : \mathbb{R} \rightarrow \{0, 1\} \mid h \text{ is any function}\}$$

This has infinite VC-dimension since we can shatter any finite set of real numbers by defining functions that assign arbitrary binary labels to those points:  $d_{VC}(\mathcal{H}) = \infty$ .

Data Distribution  $p(X, Y)$ :

$$p(X = x, Y = y) = \begin{cases} 1 & \text{for } (x, y) = (0, 0) \\ 0 & \text{for all other } (x, y) \text{ pairs} \end{cases}$$

With this construction. For any sample  $S$  of  $n > 100$  points:  $S = \{(0, 0), (0, 0), \dots, (0, 0)\}$  ( $n$  copies) and for any hypothesis  $h \in \mathcal{H}$ :

- Empirical loss:  $\hat{L}(h, S) = \frac{1}{n} \sum_{i=1}^n \ell(h(0), 0) = \ell(h(0), 0)$



- Expected loss:  $L(h) = \mathbb{E}[\ell(h(X), Y)] = \ell(h(0), 0)$

Therefore:  $L(h) = \hat{L}(h, S)$  for all  $h \in \mathcal{H}$  yielding a generalization gap of

$$L(h) - \hat{L}(h, S) = 0 \leq 0.01 \quad \text{for all } h \in \mathcal{H}$$

This holds with probability 1 (deterministically), which satisfies:

$$\mathbb{P} \left( \sup_{h \in \mathcal{H}} |L(h) - \hat{L}(h, S)| \leq 0.01 \right) = 1 \geq 0.95$$

## Appendix A

### Alternating minimization algorithm

```
def minimize_bound(L_val):
    """
    # Perform alternating minimisation of F(rho, lambda)$
    # (RHS of Eq. (13)).
    * L_val : vector (m,) of validation losses

    Returns rho, lambda^*, bound value.
    """
    m = len(L_val)
    log_pi = -np.log(m) * np.ones(m) # uniform prior \pi(h)=1/m
    (from page 10)
    lambda_ = 1.0

    for i in range(50):
        # Eq. (7): computed in log-domain for numerical stability
        log_rho = log_pi - lambda_ * n_r * L_val
        log_rho -= logsumexp(log_rho)
        rho = np.exp(log_rho)

        # helpers for Eq. (8) & bound. KL divergence is computed
        # in log-domain using the first equality in Eq. (9)
        KL = np.sum(rho * (log_rho - log_pi))
        L_emp = np.dot(rho, L_val)

        # Eq. (8): lambda_new
        num = 2.0
        denom = (
            np.sqrt(2.0 * n_r * L_emp / (KL + np.log(2 * np.sqrt(
n_r) / delta)))
            + 1.0
            + 1.0
        )
        lambda_new = num / denom

        if abs(lambda_new - lambda_) < 1e-5: # conv. criterion
            lambda_ = lambda_new
            break
    lambda_ = lambda_new
```

```

# Eq. (13): value of PAC-Bayes-\lambda bound (randomized
classifier)
bound = L_emp / (1 - lambda_ / 2.0) + (KL + np.log(2 * np.sqrt
(n_r) / delta)) / (
    n_r * lambda_ * (1 - lambda_ / 2.0)
)

return rho, lambda_, bound

```

Listing 1: Alternating minimization

## Training algorithm

```

# Build ensemble of m weak SVMs (Sec. 5 construction
# of Hypothesis Space H)
for i in range(m):
    # Select r = d + 1 random indices for training
    # and the rest for validation
    train_idx = rng.choice(n, r, replace=False)
    val_idx = np.setdiff1d(all_indices, train_idx, assume_unique=
True)

    X_tr, y_tr = X_S[train_idx], y_S[train_idx]
    X_val, y_val = X_S[val_idx], y_S[val_idx]

    # Select random gamma from grid - p. 10 in paper
    gamma = rng.choice(gamma_grid)
    clf = SVC(kernel="rbf", C=1.0, gamma=gamma).fit(X_tr, y_tr)

    L_val[i] = zero_one_loss(clf.predict(X_val), y_val)
    predsT[i] = clf.predict(X_T)

rho, _, bound = minimize_bound(L_val)

# rho-weighted Majority vote of Eq. (3.31)
# in Machine Learning: The science of selection under uncertainty
mv_raw = rho @ predsT
mv_pred = np.sign(mv_raw) # +/- 1 (np.sign(0) returns 0 - no tie
here)

```

Listing 2: Training algorithm

## References

- [Sig+89] V. Sigillito et al. *Ionosphere*. UCI Machine Learning Repository. DOI: <https://doi.org/10.24432/C5W01B>. 1989.
- [Gun99] Gregory Gundersen. *The Log-Sum-Exp Trick*. 1999. URL: <https://gregorygundersen.com/blog/2020/02/09/log-sum-exp/> (visited on 05/23/2025).

- [JDH99] Tommi Jaakkola, Mark Diekhans, and David Haussler. “Using the Fisher Kernel Method to Detect Remote Protein Homologies”. In: *Proceedings of the Seventh International Conference on Intelligent Systems for Molecular Biology*. AAAI Press, 1999, pp. 149–158. ISBN: 1577350839.
- [HCL03] Chih-Wei Hsu, Chih-Chung Chang, and Chih-Jen Lin. *A Practical Guide to Support Vector Classification*. Tech. rep. Department of Computer Science, National Taiwan University, 2003. URL: <http://www.csie.ntu.edu.tw/~cjlin/papers.html>.
- [CL11] Chih-Chung Chang and Chih-Jen Lin. “LIBSVM: A library for support vector machines”. In: *ACM Transactions on Intelligent Systems and Technology* 2 (3 2011). Software available at <http://www.csie.ntu.edu.tw/~cjlin/libsvm>, 27:1–27:27.
- [Ped+11] F. Pedregosa et al. “Scikit-learn: Machine Learning in Python”. In: *Journal of Machine Learning Research* 12 (2011), pp. 2825–2830.
- [Thi+17] Niklas Thiemann et al. *A Strongly Quasiconvex PAC-Bayesian Bound*. 2017. arXiv: 1608.05610 [cs.LG]. URL: <https://arxiv.org/abs/1608.05610>.
- [Sel25] Yevgeny Seldin. *Machine Learning The Science of Selection under Uncertainty*. Apr. 22, 2025. URL: <https://sites.google.com/site/yevgenyseldin/teaching>.